

```

[hduser@localhost ~]$ jps
4113 ResourceManager
3682 DataNode
3940 SecondaryNameNode
4235 NodeManager
3500 NameNode
6686 Jps
[hduser@localhost ~]$ hive --service metastore
ls: cannot access /usr/local/spark/lib/spark-assembly-*.jar: No such file or directory
Starting Hive Metastore Server
22/06/30 10:34:23 WARN conf.HiveConf: HiveConf of name hive.mapjoin.optimized.keys does not exist
22/06/30 10:34:23 WARN conf.HiveConf: HiveConf of name hive.optimize.multigroupby.common.distincts does not exist
22/06/30 10:34:23 WARN conf.HiveConf: HiveConf of name hive.mapjoin.lazy.hashtable does not exist
22/06/30 10:34:23 WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.min.worker.threads does not exist
22/06/30 10:34:23 WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.max.worker.threads does not exist
22/06/30 10:34:23 WARN conf.HiveConf: HiveConf of name hive.server2.logging.operation.verbose does not exist
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/phoenix-4.11.0-HBase-0.98-client.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the
SPI and manual loading of the driver class is generally unnecessary.
S

```

```

[hduser@localhost ~]$ hive
ls: cannot access /usr/local/spark/lib/spark-assembly-*.jar: No such file or directory
22/06/30 10:36:25 WARN conf.HiveConf: HiveConf of name hive.mapjoin.optimized.keys does not exist
22/06/30 10:36:25 WARN conf.HiveConf: HiveConf of name hive.optimize.multigroupby.common.distincts does not exist
22/06/30 10:36:25 WARN conf.HiveConf: HiveConf of name hive.mapjoin.lazy.hashtable does not exist
22/06/30 10:36:25 WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.min.worker.threads does not exist
22/06/30 10:36:25 WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.max.worker.threads does not exist
22/06/30 10:36:25 WARN conf.HiveConf: HiveConf of name hive.server2.logging.operation.verbose does not exist

Logging initialized using configuration in jar:file:/usr/local/hive/lib/hive-common-1.2.2.jar!/hive-log4j.properties
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/phoenix-4.11.0-HBase-0.98-client.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
hive>

```

```

hive> show databases;
OK
cognizant
cts
custdb
default
insure
retail
retailstg
sample
Time taken: 1.084 seconds, Fetched: 8 row(s)
hive> create database demo_db;
OK
Time taken: 1.75 seconds
hive> use demo_db
> ;
OK
Time taken: 0.172 seconds
hive> create table customers(
> id bigint,
> name string,
> address string);
OK
Time taken: 1.424 seconds
hive> S

```

```

hive> insert into customers values (1,"Deepti","Hyderabad"),(2,"Alekhya","Nizamabad"),(3,"Likhith","Pune"),(4,"Keerthana","Vizag");
Query ID = hduser_20220630104932_7f8da82f-deaf-48ff-b346-75adaldd0fff
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the
SPI and manual loading of the driver class is generally unnecessary.
Starting Job = job_1656561501787_0001, Tracking URL = http://Inceptez:8088/proxy/application_1656561501787_0001/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1656561501787_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2022-06-30 10:50:08,559 Stage-1 map = 0%, reduce = 0%
2022-06-30 10:50:20,241 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.72 sec
MapReduce Total cumulative CPU time: 2 seconds 720 msec
Ended Job = job_1656561501787_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://localhost:54310/user/hive/warehouse/demo_db.db/customers/.hive-staging_hive_2022-06-30_10-49-32_604_8908191457049199529-1/-ext-10000
Loading data to table demo_db.customers
[Error 30017]: Skipping stats aggregation by error org.apache.hadoop.hive.ql.metadata.HiveException: [Error 30001]: StatsPublisher cannot be initialized. The
re was a error in the initialization of StatsPublisher, and retrying might help. If you dont want the query to fail because accurate statistics could not be
collected, set hive.stats.reliable=false
Table demo_db.customers stats: [numFiles=1, totalSize=72]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 2.72 sec HDFS Read: 3868 HDFS Write: 72 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 720 msec
OK
Time taken: 50.969 seconds
hive>

```

```

hive> describe formatted customers;
OK
# col_name          data_type          comment
id                  bigint
name                string
address             string

```

```

# Detailed Table Information
Database:           demo_db
Owner:              hduser
CreateTime:         Thu Jun 30 10:43:03 IST 2022
LastAccessTime:     UNKNOWN
Protect Mode:       None
Retention:          0
Location:           hdfs://localhost:54310/user/hive/warehouse/demo_db.db/customers
Table Type:         MANAGED_TABLE
Table Parameters:
    COLUMN_STATS_ACCURATE    true
    numFiles                  1
    totalSize                 72
    transient_lastDdlTime     1656566423

# Storage Information
SerDe Library:      org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:        org.apache.hadoop.mapred.TextInputFormat
OutputFormat:       org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:         No
Num Buckets:        -1
Bucket Columns:     []
Sort Columns:       []
Storage Desc Params:
    serialization.format      1
Time taken: 0.575 seconds, Fetched: 31 row(s)

```

```

hive> select * from customers where address="Hyderabad";
OK
1      Deepti Hyderabad
Time taken: 1.441 seconds, Fetched: 1 row(s)

```

Load data from local system:

```
hive> create table if not exists orders(
  > id bigint,
  > product_id string,
  > customer_id int,
  > quantity int,
  > amount double)
  > row format delimited
  > fields terminated by ","
  > stored as textfile;
OK
Time taken: 0.277 seconds
hive> load data local inpath "/home/hduser/orders.csv" into table orders;
Loading data to table demo_db.orders
Table demo_db.orders stats: [numFiles=1, totalSize=107]
OK
Time taken: 0.848 seconds
hive> select * from orders;
OK
1232    Phone    1        2        22000.0
3412    Camera   3        3        30000.0
3468    Laptop   4        1        80000.0
1234    Watch    2        5        50000.0
```

Load data from hdfs:

```
[hduser@localhost ~]$ hdfs dfs -put orders.csv /user/hduser
22/06/30 11:39:37 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

hive> load data inpath "/user/hduser/orders.csv" into table orders;
Loading data to table demo_db.orders
Table demo_db.orders stats: [numFiles=2, totalSize=221]
OK
Time taken: 2.452 seconds
hive> select * from orders;
OK
1232    Phone    1        2        22000.0
3412    Camera   3        3        30000.0
3468    Laptop   4        1        80000.0
1234    Watch    2        5        50000.0
1457    Desktop  1        30000    NULL
1232    Phone1   1        2        22000.0
3412    Camera1  3        3        30000.0
3468    Laptop1  4        1        80000.0
1234    Watch1   2        5        50000.0
1457    Desktop1 1        1        30000.0
Time taken: 1.115 seconds, Fetched: 10 row(s)
hive> S
```